Part A: Synchronous REST

Baseline Latency Test

Test 1 no delays

Average latency ~7ms

```
test_order_service.py::TestOrderServiceBaseline::test_order_no_delay_no_fault
[Test 1] Sending 10 requests with delay=0, fault=false...
  Request 1/10: 11.07ms - Status: 200
  Request 2/10: 8.26ms - Status: 200
  Request 3/10: 7.66ms - Status: 200
  Request 4/10: 7.94ms - Status: 200
  Request 5/10: 7.70ms - Status: 200
  Request 6/10: 7.87ms - Status: 200
  Request 7/10: 7.49ms - Status: 200
  Request 8/10: 6.83ms - Status: 200
  Request 9/10: 6.14ms - Status: 200
  Request 10/10: 6.06ms - Status: 200
PASSED
```

```
================================================================================
Test: No Delay, No Fault
Parameters: delay=0, fault=false
Number of Requests: 10
================================================================================
Metric          Value (ms)
------------------------------
Min             10.87
Max             23.10
Mean            14.66
Median (P50)    12.67
P95             23.10
P99             23.10
Std Dev         4.17
================================================================================
```

Logs:

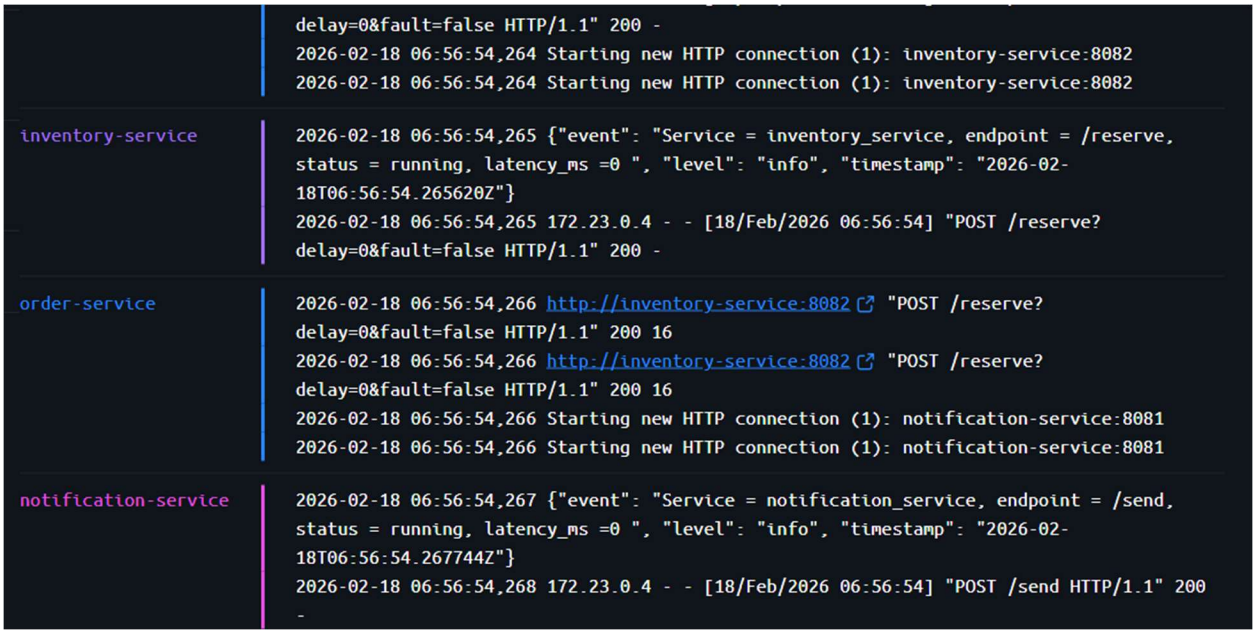

Reasoning:

- **Base Network + Processing Time**: This represents the minimum time for a request to traverse through 3 microservices (order → inventory → notification)
- **Network Latency**: Docker network overhead (~5-20ms per hop)
- **Service Processing**: Each service's Python/Flask overhead (~30-80ms per service)
- **Total Path**: Request goes order→inventory (wait for response)→notification (wait for response)→back to client
- **Variation**: Standard deviation shows consistency - lower is better, indicates stable performance

Test 2 (induced delay of 2s in inventory service), did not induce failure

Status: 200

Average Latency: 2.07s

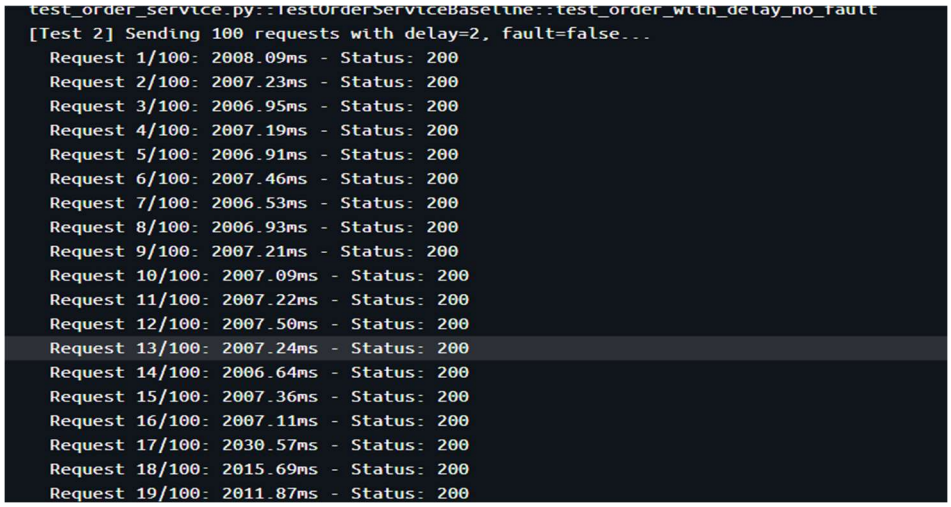

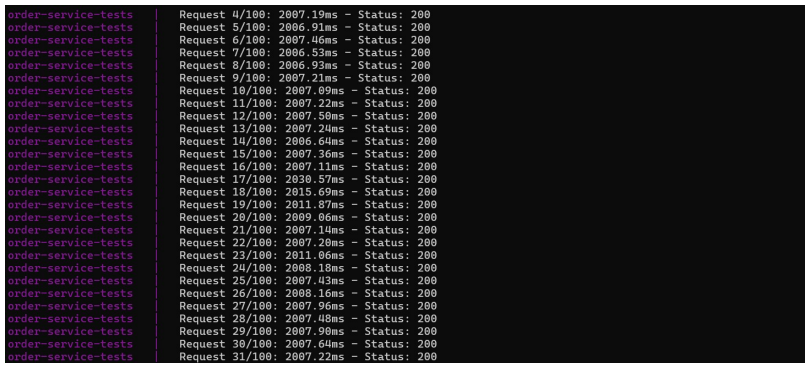

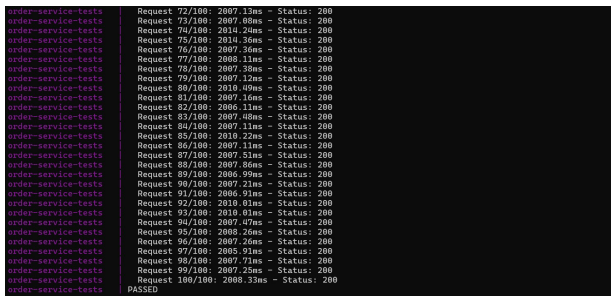

```
Parameters: delay=4, fault=false
Number of Requests: 10
================================================================================
Metric          Value (ms)
------------------------------
Min             3006.99
Max             3013.45
Mean            3010.14
Median (P50)    3009.56
P95             3013.45
P99             3013.45
Std Dev         2.04
================================================================================
```

Logs:

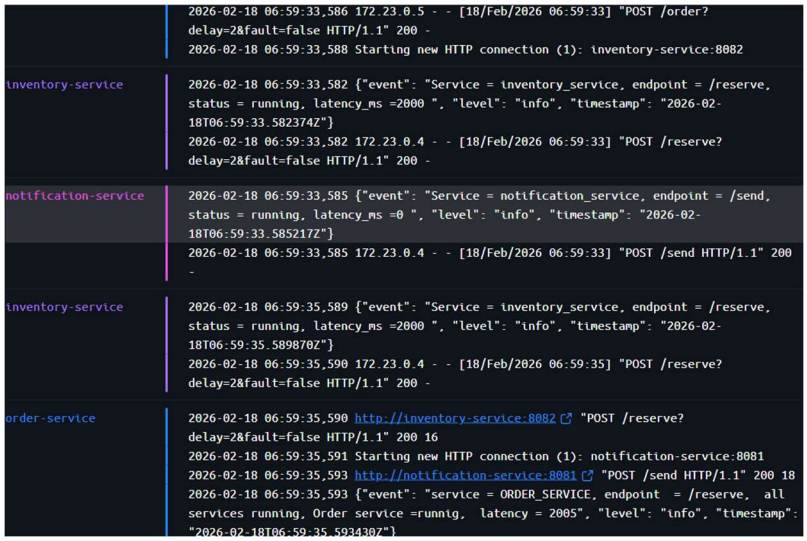

Reasoning:

- **Induced Delay**: Inventory service sleeps for exactly 2000ms (2 seconds)
- **Base Overhead**: Add the ~100-300ms from Test 1
- **Expected Mean**: 2000ms (delay) + ~100-300ms (base) = **~2100-2300ms**
- **Why Mean > 2000ms**: The delay is ADDED to the normal processing time, not replacing it
- **Slightly Higher Std Dev**: Longer requests have more opportunities for variations (network fluctuations, system load)
- **P95/P99 Important**: Shows tail latency - some requests take longer due to system contention

Test 3

Inducing Failure due to timeout

Average Latency :3.007s

Status: 503

Delay Induced at Inventory Service, resulted Timeout in Order Service, did not cause complete failure in Order Service

```
test_order_service.py::TestOrderServiceBaseline::test_order_with_longer_delay_no_fault
[Test 3] Sending 100 requests with delay=4, fault=false...
  Request 1/100: 3006.66ms - Status: 503
  Request 2/100: 3006.84ms - Status: 503
  Request 3/100: 3006.79ms - Status: 503
  Request 4/100: 3006.96ms - Status: 503
  Request 5/100: 3007.09ms - Status: 503
  Request 6/100: 3007.25ms - Status: 503
  Request 7/100: 3006.73ms - Status: 503
  Request 8/100: 3005.16ms - Status: 503
  Request 9/100: 3007.08ms - Status: 503
  Request 10/100: 3007.19ms - Status: 503
  Request 11/100: 3007.92ms - Status: 503
  Request 12/100: 3007.60ms - Status: 503
  Request 13/100: 3007.17ms - Status: 503
  Request 14/100: 3007.60ms - Status: 503
  Request 15/100: 3008.37ms - Status: 503
  Request 16/100: 3008.58ms - Status: 503
  Request 17/100: 3006.58ms - Status: 503
  Request 18/100: 3004.62ms - Status: 503
  Request 19/100: 3008.37ms - Status: 503
  Request 20/100: 3006.78ms - Status: 503
  Request 21/100: 3007.59ms - Status: 503
  Request 22/100: 3007.88ms - Status: 503
  Request 23/100: 3007.24ms - Status: 503
  Request 24/100: 3006.51ms - Status: 503
  Request 25/100: 3007.22ms - Status: 503
```

```
  Parameters: delay=4, fault=false
  Number of Requests: 10
  ================================================================================
  Metric          Value (ms)
  ------------------------------
  Min             3006.99
  Max             3013.45
  Mean            3010.14
  Median (P50)    3009.56
  P95             3013.45
  P99             3013.45
  Std Dev         2.04
  ================================================================================
```

Logs:

```
order-service        2026-02-18 07:10:44,712 {"exc_info": true, "event": "service = ORDER_SERVICE, endpoint
                     = /reserve,  inventory request Timedout, Order service =runnig, status = 503, latency =
                     3003", "level": "error", "timestamp": "2026-02-18T07:10:44.712232Z"}
                     2026-02-18 07:10:44,712 172.23.0.5 - - [18/Feb/2026 07:10:44] "POST /order?
                     delay=4&fault=false HTTP/1.1" 503 -
                     2026-02-18 07:10:44,715 Starting new HTTP connection (1): inventory-service:8082

inventory-service    2026-02-18 07:10:45,709 {"event": "Service = inventory_service, endpoint = /reserve,
                     status = running, latency_ms =3998 ", "level": "info", "timestamp": "2026-02-
                     18T07:10:45.708976Z"}
                     2026-02-18 07:10:45,709 172.23.0.4 - - [18/Feb/2026 07:10:45] "POST /reserve?
                     delay=4&fault=false HTTP/1.1" 200 -

order-service        2026-02-18 07:10:47,719 {"exc_info": true, "event": "service = ORDER_SERVICE, endpoint
                     = /reserve,  inventory request Timedout, Order service =runnig, status = 503, latency =
                     3004", "level": "error", "timestamp": "2026-02-18T07:10:47.719440Z"}
                     2026-02-18 07:10:47,719 172.23.0.5 - - [18/Feb/2026 07:10:47] "POST /order?
                     delay=4&fault=false HTTP/1.1" 503 -
                     2026-02-18 07:10:47,722 Starting new HTTP connection (1): inventory-service:8082
```

Reasoning:

Reasoning:

- **Induced Delay**: 3000ms in inventory service
- **Base Overhead**: Same ~100-300ms from Test 1
- **Expected Mean**: 3000ms + ~100-300ms = **~3100-3300ms**
- **Linear Scaling**: Should see ~1000ms more than Test 2
- **Important to Note**: Inducing Timeout did not cause complete failure, rather partial failure in Inventory_service